



Under-the-hood: Creating Your Own Spark Data Sources

Speaker: Jayesh Thakrar @ Conversant



Why Build Your Own Data Source?

Don't need

- text, csv, ORC, Parquet, JSON files
- JDBC sources
- When available from project / vendor, e.g.
 - Cassandra, Kafka, MongoDB, etc
 - Teradata, Greenplum (2018), etc

Need to

- Use special features, e.g.
 - Kafka transactions
 - Exploit RDBMS specific partition features
- Because you can, and want to ☺

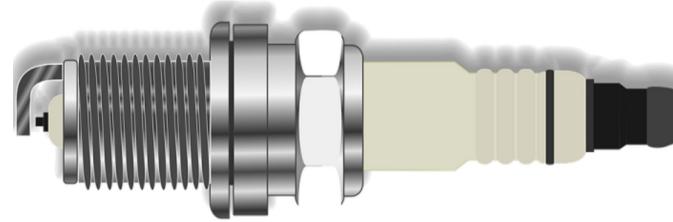
Conversant Use Cases (mid-2017)

- Greenplum as data source (read/write)
- Incompatibility between Spark, Kafka and Kafka connector versions

Agenda

1. Introduction to Spark data sources
2. Walk-through sample code
3. Practical considerations

Introduction To Spark Data Source



Using Data Sources

- **Built-in**

```
spark.read.csv("path")
```

```
spark.read.orc("path")
```

```
spark.read.parquet("path")
```

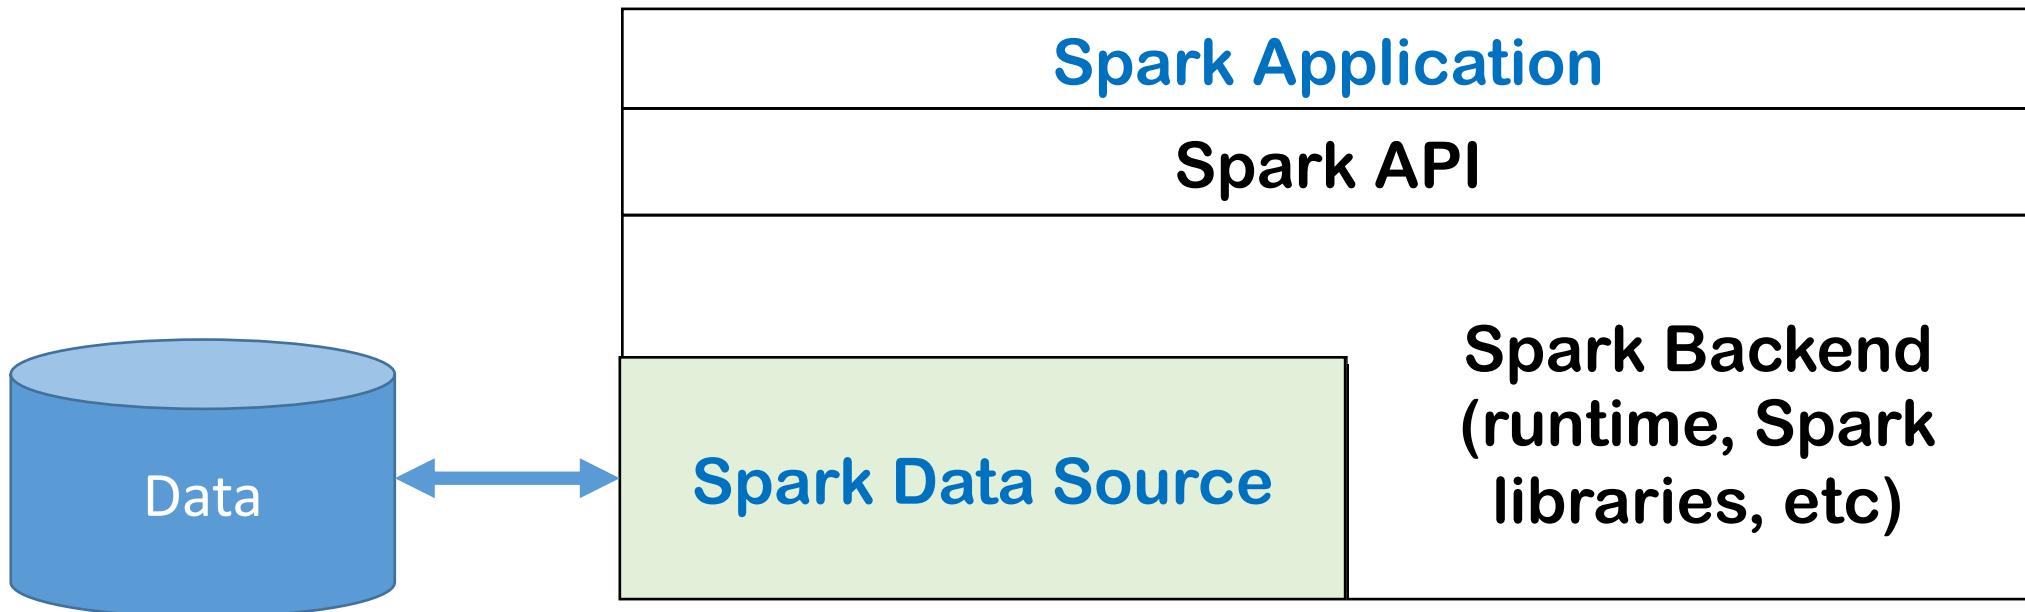
- **Third-party/custom**

```
spark.read.format("class-name").load() // custom data source
```

```
spark.read.format("...").option("...", "...").load() // with options
```

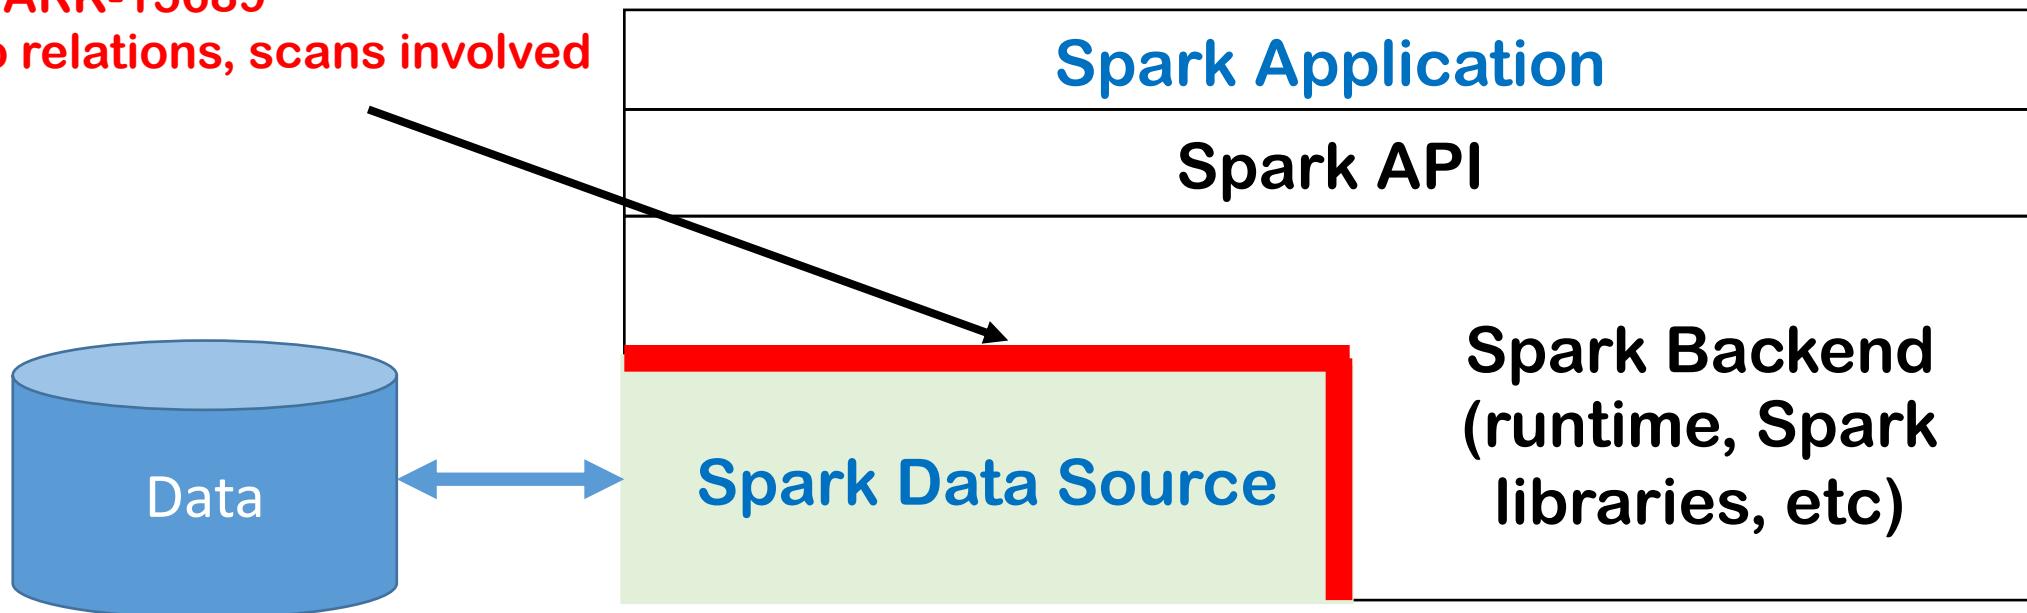
```
spark.read.format("...").schema("...").load() // with schema
```

Spark Data Source



Data Source V2 API

Shiny, New V2 API
since Spark 2.3 (Feb 2018)
SPARK-15689
No relations, scans involved



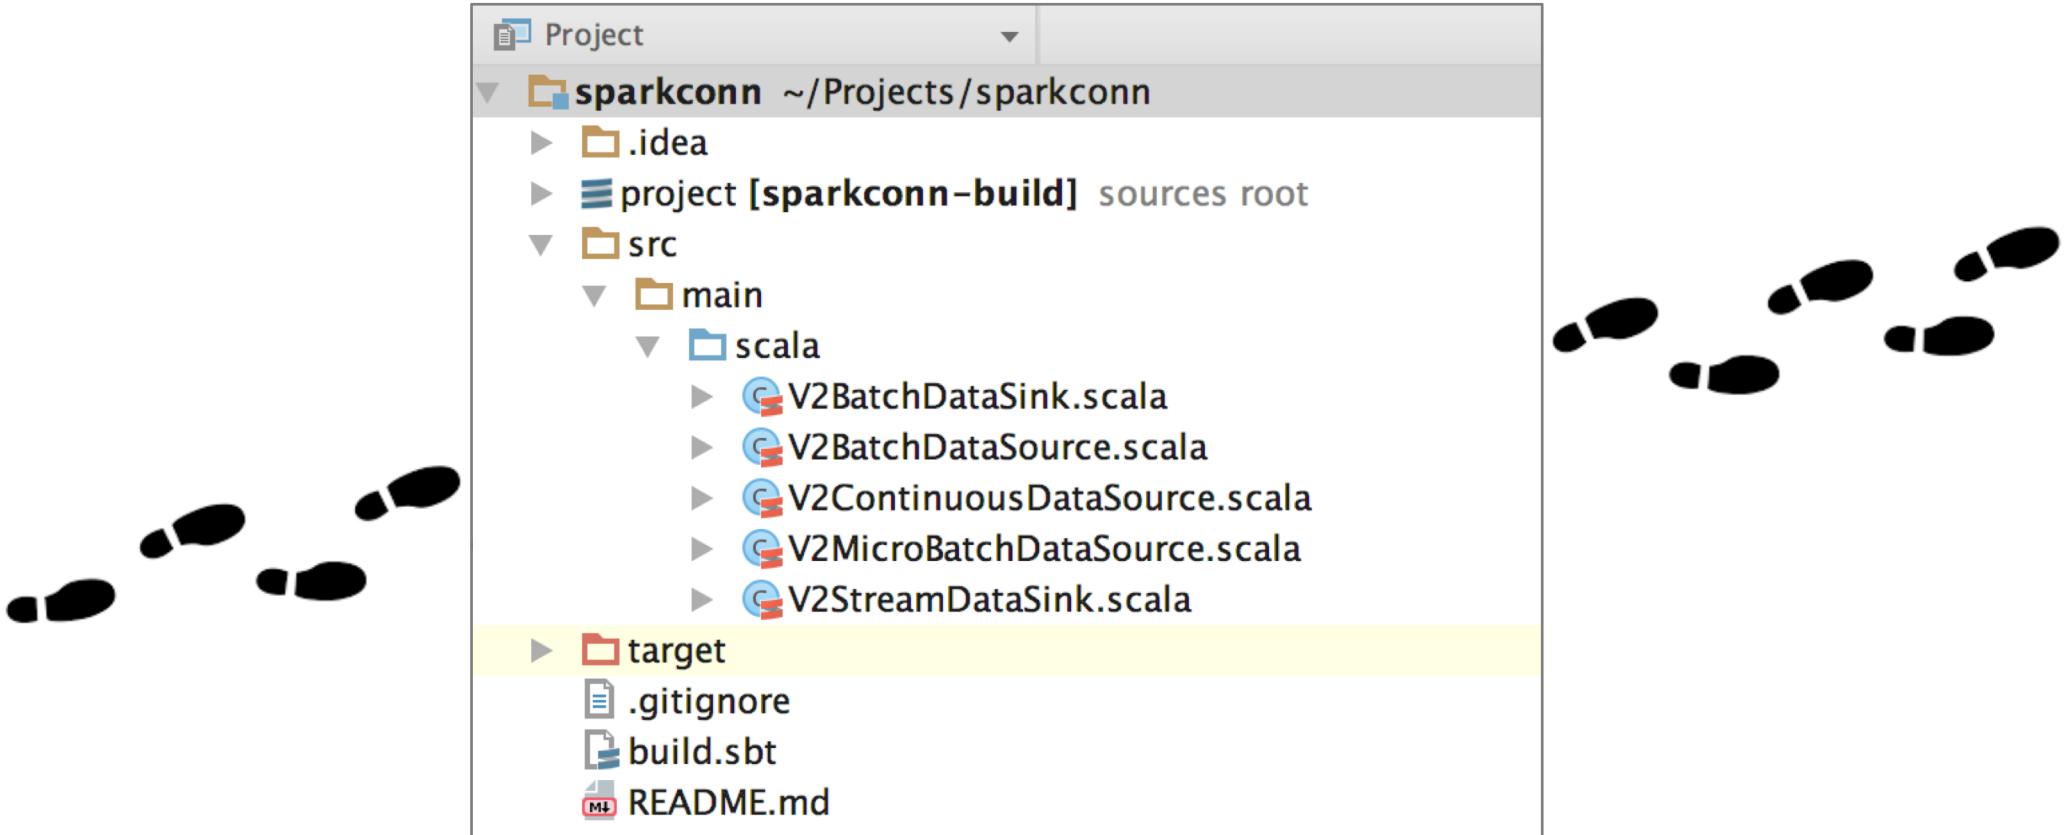
V2 API: Data Source Types

Data Source Type	Output
Batch	Dataset
Microbatch (successor to Dstreams)	Structured Stream = stream of bounded dataset
Continuous	Continuous Stream = continuous stream of Row(s)

V2 DATA SOURCE API

- **Well documented design, but still evolving**
V2 API : [Design Doc](#), Feature SPARK-15689
Continuous Streaming Design Doc and Feature - SPARK-20928
- **Implemented as Java Interfaces (not classes)**
- **Similar interfaces across all data source types**
- **Microbatch and Continuous not hardened yet...**
e.g. SPARK-23886, SPARK-23887, SPARK-22911

Code Walkthrough



Details

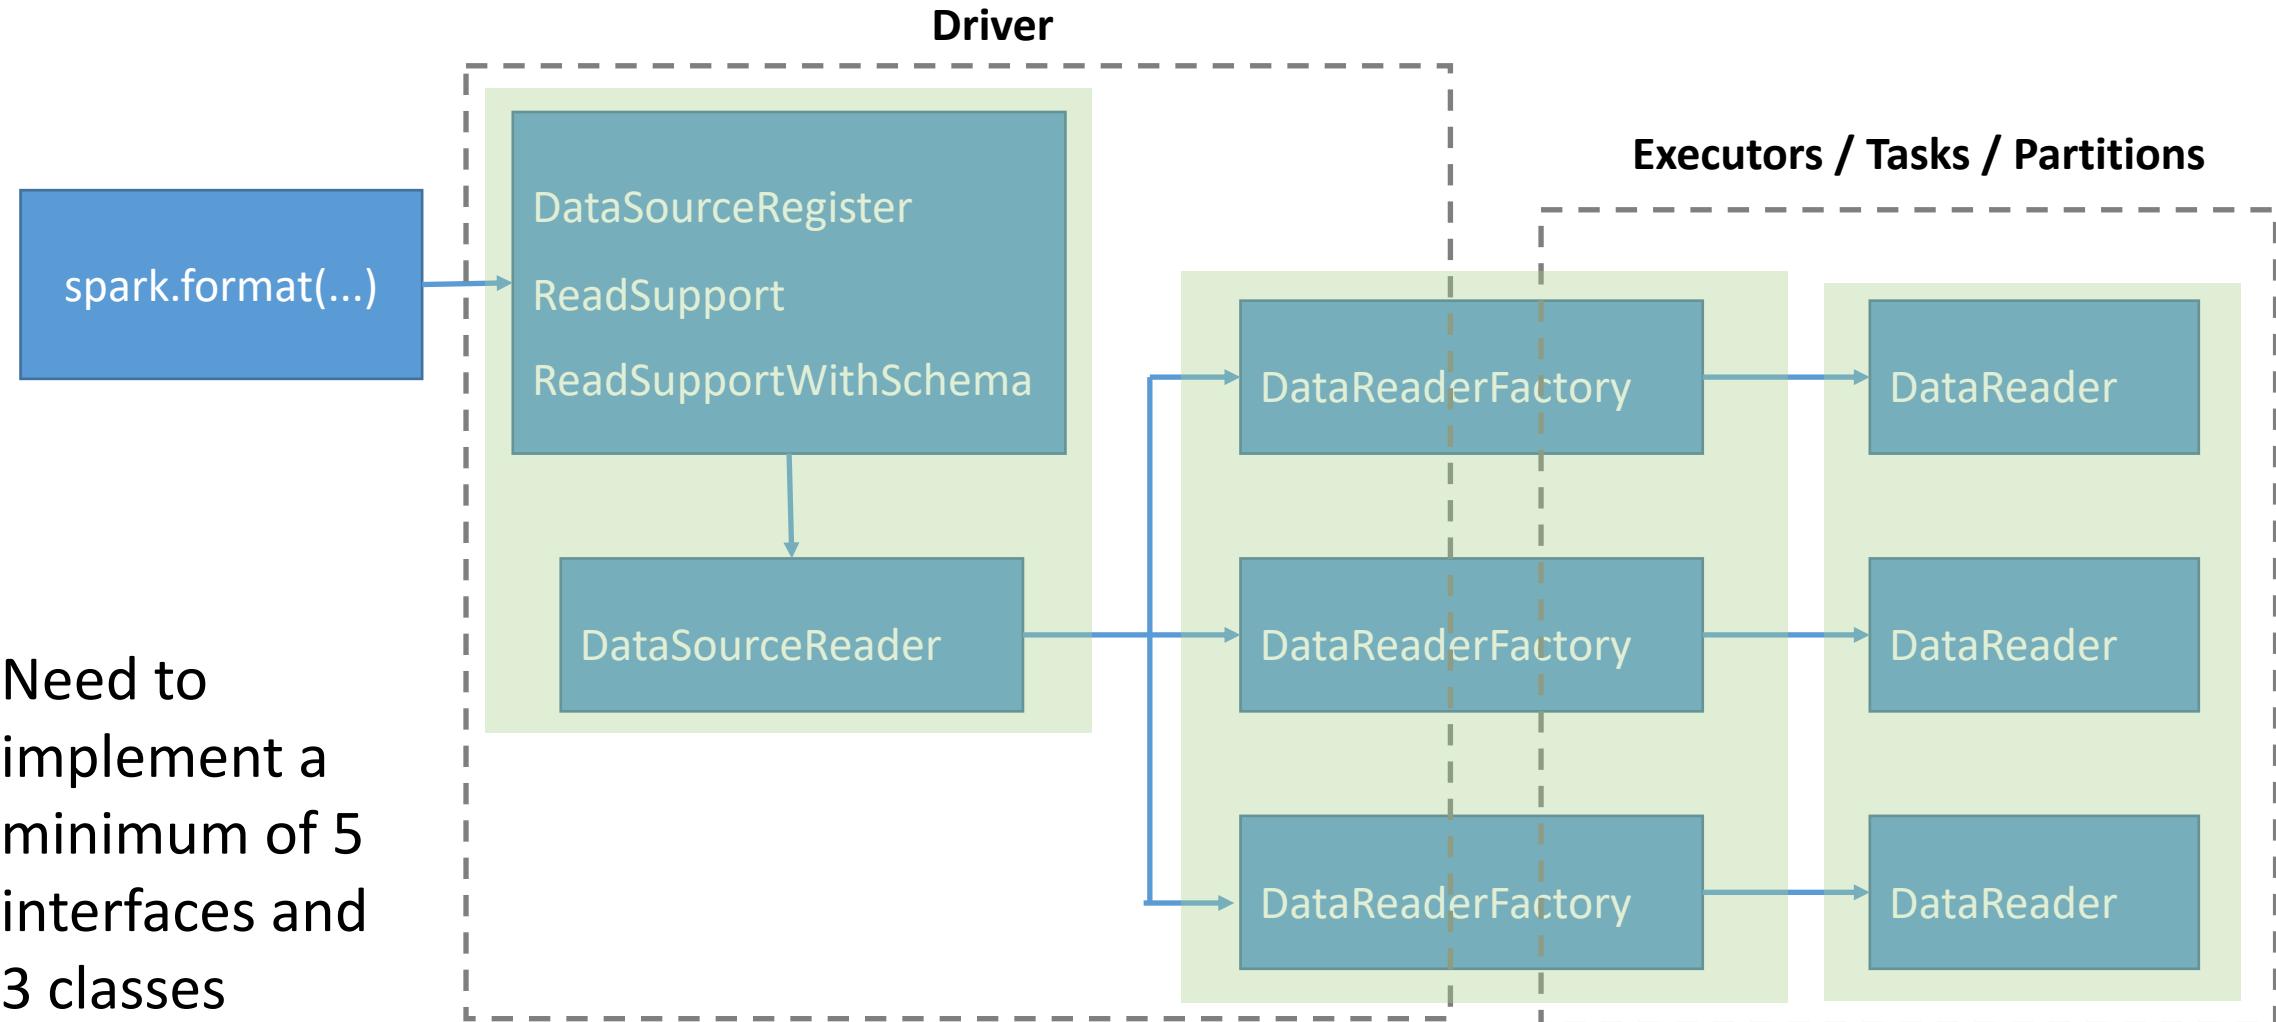
- Project: <https://github.com/JThakrar/sparkconn>
- Requirements:
 - Scala 2.11
 - SBT

Reading From Data Source

```
val data = spark.read.format("DataSource").option("key", "value").schema(..).load()
```

Step	Action
spark	= SparkSession
read	Returns a DataFrameReader for data source orchestration
format	Lazy lookup of data source class - by shortname or by full-qualified class name
option	Zero, one or more key-value pairs of options for data source
schema	Optional, user-provided schema
load	Loads the data as a dataframe. Remember dataframes are lazy-evaluated.

Reading: Interfaces to Implement



Interface Definitions And Dependency



```
libraryDependencies +=  
  "org.apache.spark" %%  
  "spark-sql" %  
  "2.3.1" %  
  "provided"
```

Example V2 API Based Datasource

- Very simple data source that generates a row with a fixed schema of a single string column. Column name = "string_value"
- Completely self-contained (i.e. no external connection)
- Number of partitions in dataset is user-configurable (default = 5)
- All partitions contains same number of rows (strings)
- Number of rows per partition is user-configurable (default = 5)

Read Interface: DataSourceRegister

Interface	Purpose
org.apache.spark.sql.sources.v2.DataSourceRegister	DataSourceRegister is the entry point for your data source.
org.apache.spark.sql.sources.v2.ReadSupport and / or org.apache.spark.sql.sources.v2.ReadSupportWithSchema	ReadSupport is then responsible to instantiate the object implementing DataSourceReader (discussed later). It accepts the options/parameters and optional schema from Spark application.

```
class V2BatchDataSource
  extends DataSourceRegister
  with ReadSupport { // use ReadSupportWithSchema to support user-provided Schema

  val DEFAULT_PARTITION_COUNT: Int = 5

  val DEFAULT_ROWS_PER_PARTITION: Int = 5

  override def shortName(): String = "v2batchdatasource" // to implement DataSourceRegister

  override def createReader(options: DataSourceOptions): DataSourceReader = { // to implement ReadSupport
    val optionsKV = options.asMap().asScala // converts options to lower-cased keyed map
    val partitions = optionsKV.getOrDefault("partitions", DEFAULT_PARTITION_COUNT.toString).toInt
    val rows = optionsKV.getOrDefault("rowsperpartition", DEFAULT_ROWS_PER_PARTITION.toString).toInt
    assert(partitions > 0, s"Partitions should be > 0 (specified value = $partitions)")
    assert(rows > 0, s"Rows should be > 0 (specified value = $rows)")
    println(s"\n\nCreating ${this.getClass.getName}: with $partitions partitions, each with $rows rowsPerPartition\n")
    new V2BatchDataSourceReader(partitions, rows)
  }
}
```

Read Interface: DataSourceReader

Interface	Purpose
org.apache.spark.sql.sources.v2.reader.DataSourceReader	This interface requires implementations for: <ul style="list-style-type: none">• determining the schema of data• determining the number of partitions and creating that many reader factories below.

```
class V2BatchDataSourceReader(partitions: Int,
                             rows: Int)
  extends DataSourceReader {

  println(s"\n\nCreating ${this.getClass.getName}: $partitions $partitions and $rows $rowsPerPartition each\n")

  override def readSchema(): StructType = {
    StructType(StructField("string_value", StringType, false)::Nil)
  }

  override def createDataReaderFactories(): java.util.List[DataReaderFactory[Row]] = {
    println("Creating DataReaderFactories.....")
    (1 to partitions).map(partition =>
      new V2BatchDataReaderFactory(partition, rows, partitions).asInstanceOf[DataReaderFactory[Row]].asJava
    )
  }
}
```

Read Interface: DataReaderFactory

Interface	Purpose
org.apache.spark.sql.sources.v2.reader. DataReaderFactory	This is the "handle" passed by the driver to each executor. It instantiates the readers below and controls data fetch.

```
class V2BatchDataReaderFactory(partition: Int,
                               rows: Int,
                               totalPartitions: Int)
  extends DataReaderFactory[Row] {

  println(s"\n\nCreating ${this.getClass.getName}: $partition of $totalPartitions\n")

  override def createDataReader(): DataReader[Row] = {
    new V2BatchDataReader(partition, rows)
  }
}
```

Read Interface: DataReader

Interface	Purpose
org.apache.spark.sql.sources.v2.reader. DataReader	This does the actual work of fetching data from source (at task-level)

```
class V2BatchDataReader(partition: Int,
                       rows: Int)
extends DataReader[Row] {

private var rowsRemaining = rows

println(s"\n\nCreating ${this.getClass.getName}: $partition\n")

override def next(): Boolean = rowsRemaining > 0

override def get(): Row = {
  var resultRow: Row = Row("")

  if (next) {
    rowsRemaining = rowsRemaining - 1
    resultRow = Row(s"Partition: $partition || Row ${rows - rowsRemaining} of $rows")
  } else {
    new IllegalArgumentException
  }
  resultRow
}

override def close(): Unit = {}
}
```

Summary

Because you are implementing interfaces, **YOU** can determine the "class" parameters and initialization

Interface	Purpose
org.apache.spark.sql.sources.v2. DataSourceRegister org.apache.spark.sql.sources.v2. ReadSupport	DataSourceRegister is the entry point for your connector. ReadSupport is then responsible to instantiate the object implementing DataSourceReader below. It accepts the options/parameters and optional schema from Spark application.
org.apache.spark.sql.sources.v2.reader. DataSourceReader	This interface requires implementations for: <ul style="list-style-type: none">• determining the schema of data• determining the number of partitions and creating that many reader factories below. The DataSourceRegister , DataSourceReader and DataReaderFactory are instantiated at the driver. The driver then serializes DataReaderFactory and sends it to each of the executors.
org.apache.spark.sql.sources.v2.reader. DataReaderFactory	This is the "handle" passed by the driver to each executor. It instantiates the readers below and controls data fetch.
org.apache.spark.sql.sources.v2.reader. DataReader	This does the actual work of fetching data

Write Interfaces to Implement

Interface	Purpose
org.apache.spark.sql.sources.v2. DataSourceRegister org.apache.spark.sql.sources.v2. WriteSupport	DataSourceRegister is the entry point for your connector. WriteSupport is then responsible to instantiate the object implementing DataSourceWriter below. It accepts the options/parameters and the schema.
org.apache.spark.sql.sources.v2.writer. DataSourceWriter	This interface requires implementations for: <ul style="list-style-type: none">• committing data write• aborting data write• creating writer factories The DataSourceRegister , DataSourceWriter and DataWriterFactory are instantiated at the driver. The driver then serializes DataWriterFactory and sends it to each of the executors.
org.apache.spark.sql.sources.v2.writer. DataWriterFactory	This is the "handle" passed by the driver to each executor. It instantiates the readers below and controls data fetch.
org.apache.spark.sql.sources.v2.writer. DataWriter	This does the actual work of writing and committing/aborting the data
org.apache.spark.sql.sources.v2.writer. WriterCommitMessage	This is a "commit" message that is passed from the DataWriter to DataSourceWriter .

Practical Considerations



Know Your Data Source

- Configuration
- Partitions
- Data schema
- Parallelism approach
- Batch and/or streaming
- Restart / recovery

V2 API IS STILL EVOLVING

- **SPARK-22386** - Data Source V2 Improvements
- **SPARK-23507** - Migrate existing data sources
- **SPARK-24073** - DataReaderFactory Renamed in 2.4
- **SPARK-24252, SPARK-25006** - DataSourceV2: Add catalog support
- **So Why use V2?**
Future-ready and alternative to V2 needs significantly more time and effort!
See <https://www.youtube.com/watch?v=O9kpduk5D48>

About...

Conversant

- Digital marketing unit of Epsilon under Alliance Data Systems (ADS)
- (Significant) player in internet advertising.
 - We see about 80% of internet ad bids in the US
- Secret sauce = anonymous cross-device profiles driving personalized messaging

Me (Jayesh Thakrar)

- Sr. Software Engineer (jthakrar@conversantmedia.com)
-  <https://www.linkedin.com/in/jayeshthakrar/>

Questions?

