

ODSC East • Boston • 2018

May 1st – 4th • Boston Convention Center

Under-the-hood: Creating Your Own Spark Data Sources



About...

CONVERSANT

- Part of the Epsilon organization under Alliance Data
- Player in the programmatic ad bidding market (AdTech)
- Personalized messaging using non-PII anonymous profiles
(PII = Personally Identifiable Information)
- EU General Data Protection Regulation (GDPR) ready

Myself (Jayesh Thakrar)

- Sr. Software Engineer (jthakrar@conversantmedia.com)
-  <https://www.linkedin.com/in/jayeshthakrar/>

More About CONVERSANT...

The image shows the homepage of the Conversant website. At the top, there is a navigation bar with links: WHAT WE DO, WHAT WE OFFER, WHAT WE KNOW, WHO WE ARE, PUBLISHERS, CAREERS, and CONTACT. There is also a search icon. Below the navigation, a large video player displays a scene of people walking in a city street. Overlaid on the video are several text elements: "GET TO THE HEART OF WHAT MATTERS TO MILLIONS" in large, bold, white and red text; a "WATCH OUR STORY" button with a play icon; and five key statistics in boxes: "200M+ REAL PEOPLE YOU CAN MESSAGE RIGHT NOW", "7,000+ DIMENSIONS TO CONSUMER PROFILES", "96% ACCURACY AT MATCHING CONSUMERS TO DEVICES", "80% CONSUMERS STAY REACHABLE AFTER 1 YEAR", and "10X AVERAGE INCREMENTAL RETURN ON AD SPEND".

WHAT WE DO WHAT WE OFFER WHAT WE KNOW WHO WE ARE PUBLISHERS CAREERS CONTACT

GET TO THE HEART
OF WHAT MATTERS
TO MILLIONS

WATCH OUR STORY

200M+
REAL PEOPLE YOU CAN
MESSAGE RIGHT NOW

7,000+
DIMENSIONS TO
CONSUMER PROFILES

96%
ACCURACY AT MATCHING
CONSUMERS TO DEVICES

80%
CONSUMERS STAY
REACHABLE AFTER 1 YEAR

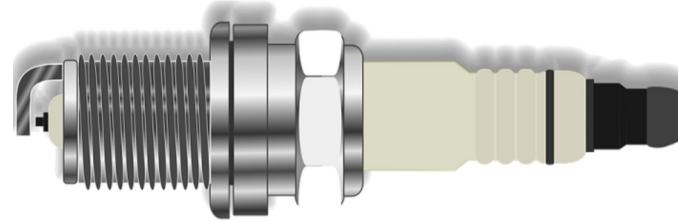
10X
AVERAGE INCREMENTAL
RETURN ON AD SPEND

Workshop Agenda

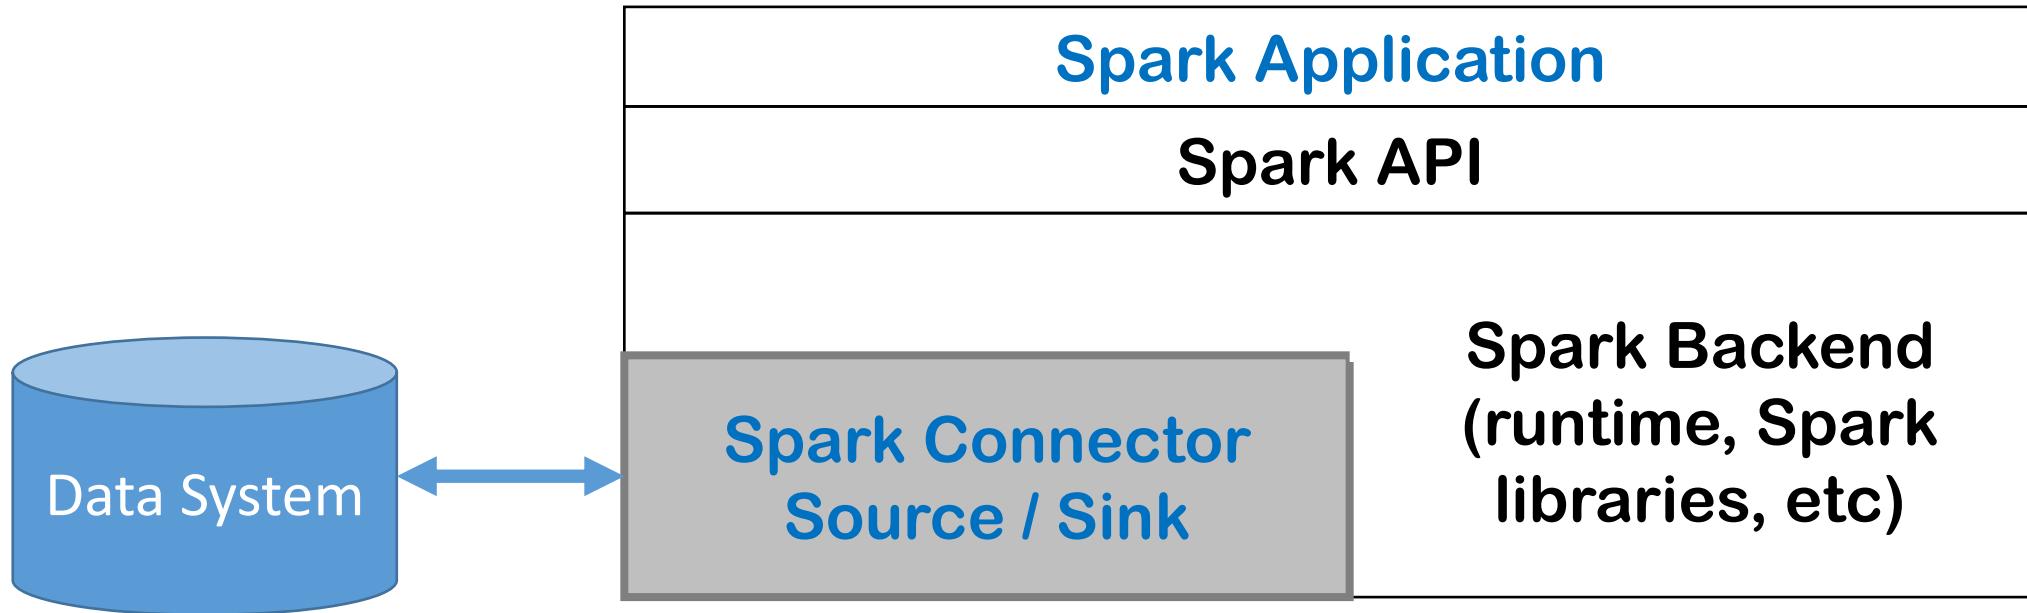
1. Introduction to Spark data sources
2. Walk-through sample code
3. Practical considerations

Slides and code on Github
<https://github.com/JThakrar/sparkconn>

Introduction To Spark Data Source

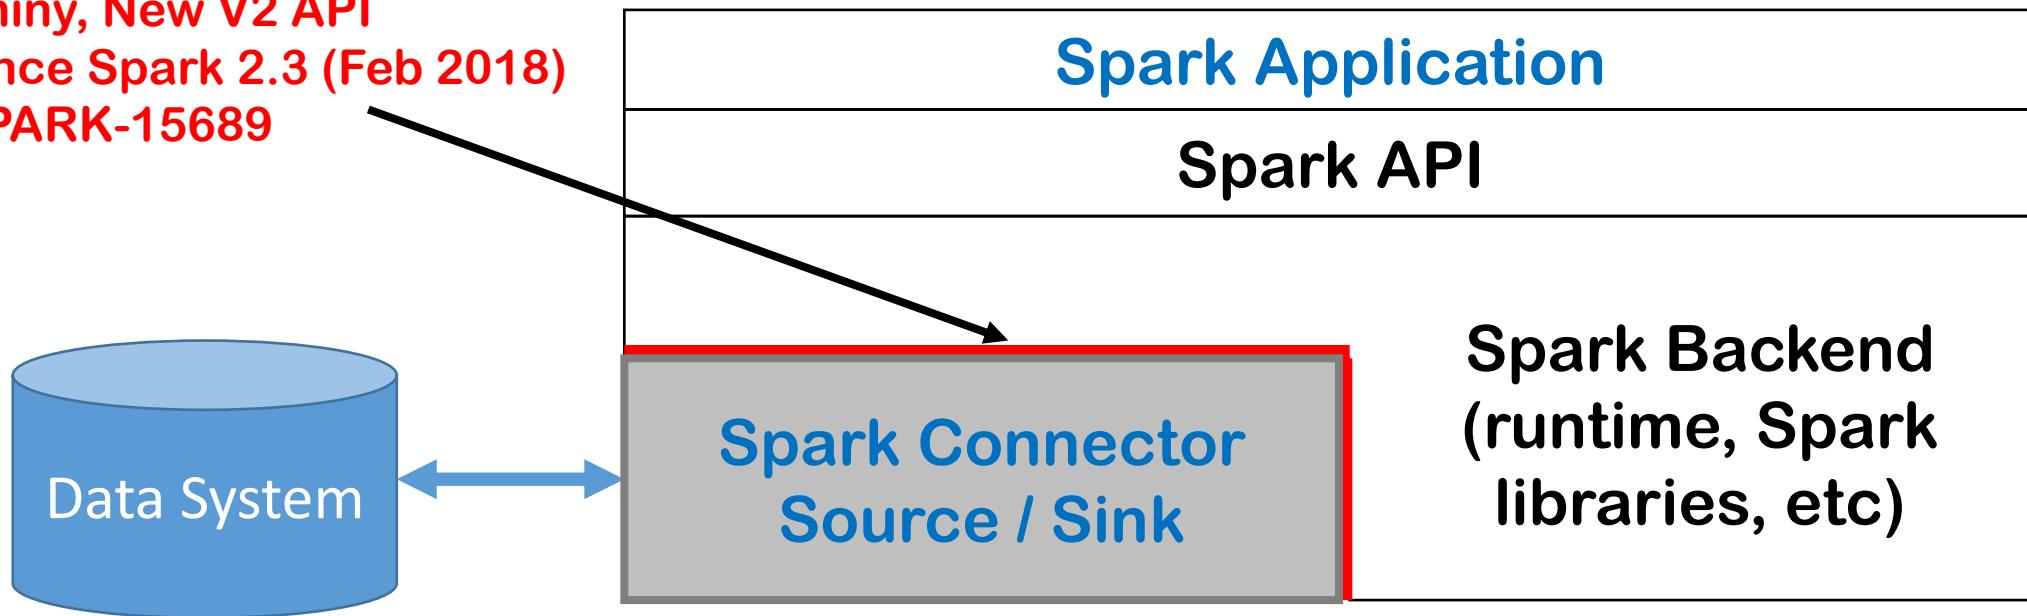


Spark Data Source



Data Source V2 API

Shiny, New V2 API
since Spark 2.3 (Feb 2018)
SPARK-15689



V2 API: Data Source Types

DataSource Type	Output
Batch	Dataset
Microbatch (successor to Dstreams)	Structured Stream = stream of bounded dataset
Continuous	Continuous Stream = continuous stream of Row(s)

V2 API

- **API = Java Interfaces**
- **Similar interfaces across all 3 data source types**
- **Write-ahead, recovery, checkpointing not ready**
- **Well-documented but still evolving**
SPARK-20928, SPARK-22386

Reading From Data Source

```
val data = spark.read.format("DataSource").option("key", "value").schema(...).load()  
data.show(100, false)
```

Step	Action
read	Instantiates a DataFrameReader for orchestration
format	Lazy lookup of data source class
schema	Optional, user-provided schema
load	Initiates actual setup of data source and instantiates the "connector"
show	Triggers actual data retrieval ("show" is an action)

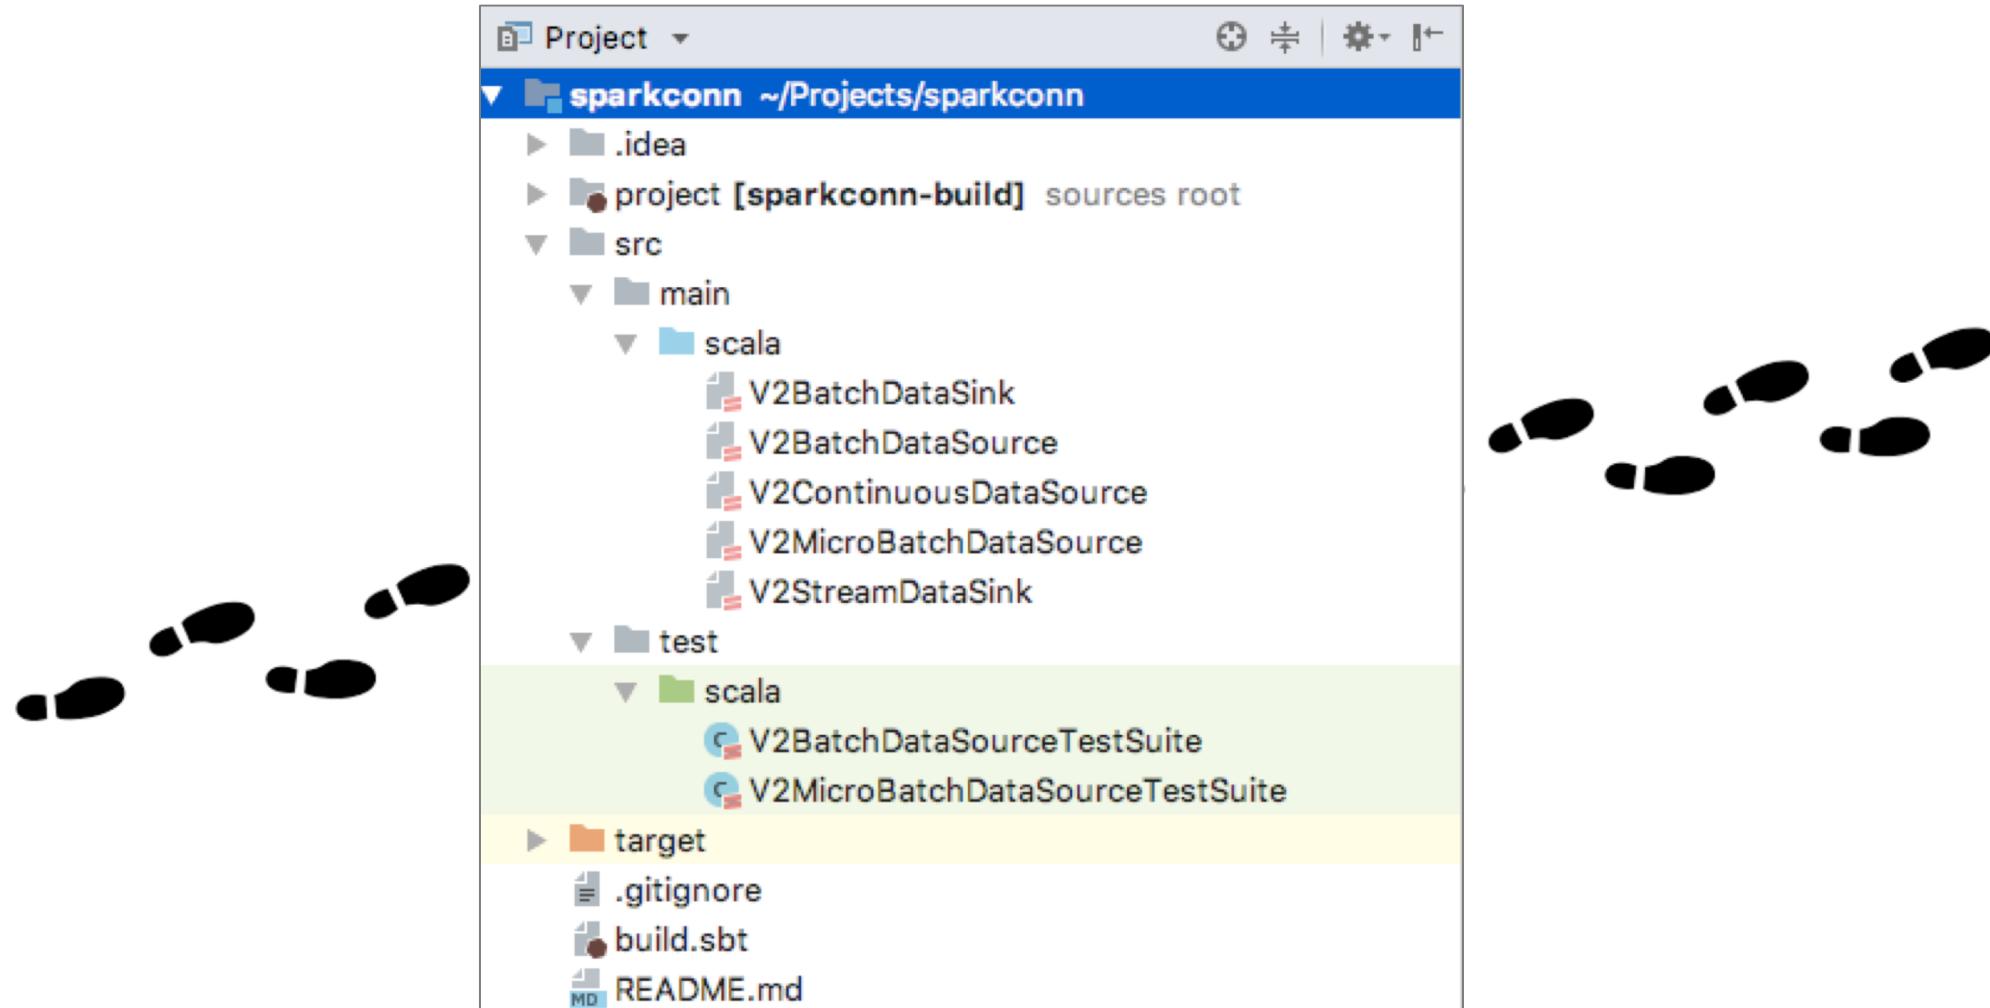
Reading: Interfaces to Implement

Java Interface	Purpose
org.apache.spark.sql.sources.v2.DataSourceRegister org.apache.spark.sql.sources.v2. ReadSupport	DataSourceRegister is the entry point for your connector. ReadSupport instantiates DataSourceReader below. Accepts options/parameters and optional schema from Spark application.
org.apache.spark.sql.sources.v2.reader. DataSourceReader	DataSourceReader needs to: <ul style="list-style-type: none">• Provide the data schema• determine the number of partitions and creating that many reader factories below.
org.apache.spark.sql.sources.v2.reader. DataReaderFactory	This is the "handle" passed by the driver to each executor. It instantiates the readers below and controls data fetch. The DataSourceRegister , DataSourceReader and DataReaderFactory are instantiated at the driver. The driver then serializes DataReaderFactory and sends it to each of the executors.
org.apache.spark.sql.sources.v2.reader. DataReader	This does the actual work of fetching data

Writing: Interfaces to Implement

Java Interface	Purpose
org.apache.spark.sql.sources.v2.DataSourceRegister org.apache.spark.sql.sources.v2. WriteSupport	DataSourceRegister is the entry point for your connector. WriteSupport instantiates DataSourceWriter . Accepts options, savemode and schema.
org.apache.spark.sql.sources.v2.writer. DataSourceWriter	This interface needs to: <ul style="list-style-type: none">register commits and aborts across all executorsdetermine the number of partitions and creating that many reader factories below.
org.apache.spark.sql.sources.v2.writer. DataWriterFactory	This is the "handle" passed by the driver to each executor. It instantiates the writers below that control the write operations. The DataSourceRegister , DataSourceWriter and DataWriterFactory are instantiated at the driver. The driver then serializes DataWriterFactory and sends it to each of the executors.
org.apache.spark.sql.sources.v2.writer. DataWriter	Does writing data as well as commits and aborts
org.apache.spark.sql.sources.v2.writer. WriterCommitMessage	Serializable empty interface to pass commit message from DataWriter to DataSourceWriter

Code Walkthrough



Details

- Project: <https://github.com/JThakrar/sparkconn>
- Requirements:
 - Spark 2.3 (<http://spark.apache.org/downloads.html>)
 - Scala 2.11 (<https://www.scala-lang.org/download/2.11.8.html>)
 - SBT (<https://www.scala-sbt.org/download.html>)
 - Editor or IDE (IntelliJ - <https://www.jetbrains.com/idea/download/>)

Practical Considerations



Know Your Data Source

- Configuration
- Partitions
- Data schema
- Parallelism approach
- Batch and/or streaming
- Restart / recovery

V2 API IS STILL EVOLVING

- **SPARK-20928**
 - SPIP: Continuous Processing Mode for Structured Streaming
- **Why use V2?**

Alternative to V2 needs significantly more time and effort!
See <https://www.youtube.com/watch?v=O9kpduk5D48>
- **Checkpointing, restartability, writeahead, etc**
- **Resource - dev@spark.apache.org**

Questions?

